Unsupervised Learning

5 试题

1 point

1. For which of the following tasks might K-means clustering be a suitable algorithm? Select all that apply.

- [ ] Given a database of information about your users, automatically group them into different market segments.
- [ ] Given sales data from a large number of products in a supermarket, figure out which products tend to form coherent groups (say are frequently purchased together) and thus should be put on the same shelf.
- [ ] Given historical weather records, predict the amount of rainfall tomorrow (this would be a real-valued output)
- [ ] Given sales data from a large number of products in a supermarket, estimate future sales for each of these products.

1 point

2. Suppose we have three cluster centroids $\mu_1 = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$, $\mu_2 = \begin{bmatrix} -3 \\ 0 \end{bmatrix}$ and $\mu_3 = \begin{bmatrix} 4 \\ 2 \end{bmatrix}$. Furthermore, we have a training example $x^{(i)} = \begin{bmatrix} -1 \\ 2 \end{bmatrix}$. After a cluster assignment step, what will $c^{(i)}$ be?

- (•) $c^{(i)} = 1$
- ( ) $c^{(i)} = 3$
- ( ) $c^{(i)} = 2$
- ( ) $c^{(i)}$ is not assigned

1 point

3. K-means is an iterative algorithm, and two of the following steps are repeatedly carried out in its inner-loop. Which two?

- [ ] Move each cluster centroid μ_k, by setting it to be equal to the closest training example $x^{(i)}$
- [ ] Move the cluster centroids, where the centroids μ_k are updated.
- [ ] The cluster assignment step, where the parameters $c^{(i)}$ are updated.
- [ ] The cluster centroid assignment step, where each cluster centroid μ_i is assigned (by setting $c^{(i)}$) to the closest training example $x^{(i)}$.

1 point

4. Suppose you have an unlabeled dataset $\{x^{(1)}, \ldots, x^{(m)}\}$. You run K-means with 50 different random

initializations, and obtain 50 different clusterings of the

data. What is the recommended way for choosing which one of

these 50 clusterings to use?

- ( ) The only way to do so is if we also have labels $y^{(i)}$ for our data.
- (•) For each of the clusterings, compute $\frac{1}{m} \sum_{i=1}^{m} ||x^{(i)} - \mu_{c(i)}||^2$, and pick the one that minimizes this.
- ( ) Always pick the final (50th) clustering found, since by that time it is more likely to have converged to a good solution.
- ( ) The answer is ambiguous, and there is no good way of choosing.

1 point

5. Which of the following statements are true? Select all that apply.

- [ ] If we are worried about K-means getting stuck in bad local optima, one way to ameliorate (reduce) this problem is if we try using multiple random initializations.
- [ ] Since K-Means is an unsupervised learning algorithm, it cannot overfit the data, and thus it is always better to have as large a number of clusters as is computationally feasible.
- [ ] The standard way of initializing K-means is setting $\mu_1 = \cdots = \mu_k$ to be equal to a vector of zeros.
- [ ] For some datasets, the "right" or "correct" value of K (the number of clusters) can be ambiguous, and hard even for a human expert looking carefully at the data to decide.

提交测试